ARDA TECHNOLOGIES

Tone of Voice

How Arda sounds in writing — across web, sales, exec voice, and customer comms.

Version **v1 — May 2026**

Prepared by **Simon Bowker, Hepburn Advisory**

For **Erik (CEO), Paul (CRO), Alex, Stefan, Costin — and anyone writing on Arda's behalf**

Status Working draft. Built on Phase 1 GTM synthesis (27 April 2026). Reviewed against ardacomputing.com, the Internal POV deck, the First Call Deck, and the Executive Slides.

How to use this document

This is a layered tone-of-voice document. It does three jobs in one:

1. **Internal reference.** Anyone at Arda writing on the company's behalf — Erik, Paul, Alex, Stefan, Costin, future hires — should be able to open this, find the relevant section, and write copy that sounds like one company.
2. **Working brief for web copy.** The principles, lexicon, and sample copy here feed directly into the next rewrite of ardacomputing.com. Section 8 (Channel Guidance — Web) and Section 11 (Before/After) are the operative parts for that work.
3. **Brand artefact.** The structure can be lifted into a designed deliverable for a copywriter, agency, or partner — the verbal half of a brand book. When that happens, sections 4 to 7 are the core.

How to read it

Sections **1 to 3** set the strategic ground (the buyer, the POV, what the voice has to do).

Sections **4 to 7** are the voice itself — DNA, principles, attributes, lexicon.

Sections **8 to 11** show how the voice plays out across channels, against reference brands, and in real copy. These are the most-used pages.

Sections **12 to 14** are the operating instructions — quick-reference card, open questions, hand-off into the website rebuild.

> **EDITORIAL POSITION**
>
> This is a working v1, not a final brand bible. The Phase 1 strategy is locked enough to commit on voice now — the alternative is publishing more copy that pulls the brand in different directions every month. Refine on contact with customers, not in committee.

1. Strategic foundations

The voice has to serve the GTM, not the product. Five facts from the Phase 1 workshop synthesis (27 April 2026) anchor every decision in this document.

Fact	What it forces in the voice	Source
Wedge bet: lead with compliance, cost secondary.	The voice leads with risk, control, governance. Cost is reinforcement, never the headline. Productivity is a side-effect, not the promise.	Synthesis §1
Lead buyer: CIO of a regulated organisation. Champion: business CXO / Head of AI. Gatekeeper: CISO / CRO.	We write to a CIO reading a board paper, not an engineer reading a spec. Plain English first; engineering depth available on request.	Synthesis §2
Point of view: Arda is the AI customers Own, Operate and Govern within their existing trust boundary.	Three verbs (Own, Operate, Govern) become structural. Sovereignty = control, not geography. We never let the language drift back to location.	Synthesis §3
Primary segment: Australian energy networks. Secondary: financial services.	Concrete worlds — substations, dispatch, claims, underwriting. Not 'enterprise' in the abstract. We name the workload.	Synthesis §4
Open: on-prem vs hyperscaler-only deployment is unresolved.	Until it lands, the voice talks about deployment as flexible (private cloud, on-prem, hybrid). We do not pick a fight in copy that leadership has not picked internally.	Synthesis §5 (Open items)

2. The job of the voice

Before we can describe what the voice sounds like, we have to be clear what work it has to do. The CIO of an Australian transmission network reading our homepage on a Tuesday morning has three thoughts in this order:

4. Will this end my career? (Risk)
5. Can I defend this to my CISO and my board? (Authority)
6. Does the business case stack up? (Cost / value)

The voice has to answer those three questions in that order. Not the founder's order — the buyer's order. The current website answers them in the wrong order: it leads with 'infrastructure layer' (engineering pride), productivity uplift (cost), then security as a feature list. That works for an engineering audience. It does not work for a CIO.

> **THE SINGLE TEST**
> Could a CIO read this aloud in their next risk committee without feeling exposed? If yes, it's on-voice. If no, rewrite.

What the voice is NOT for

- It is not a developer-relations voice. We do not perform engineering depth as a marketing strategy. (That voice belongs in the technical appendix, the docs, and Stefan's blog posts — not the homepage.)
- It is not a startup-cool voice. No emojis in customer comms, no 'we're shipping' theatre, no manifestos with one-word lines.
- It is not an AI-vendor voice. We do not promise '10x productivity', 'AGI', 'transformation'. Those words are now red flags to a CIO, not green ones.

3. Voice DNA

If everything else in this document is lost, this is the line to keep:

> THE DNA
>
> **Sober, sovereign, specific.** *We sound like the trusted operator who has thought about the failure modes — not the vendor selling around them.*

Three words, unpacked

Word	What it means in writing	Failure mode if we lose it
SOBER	Declarative sentences. Verbs over adjectives. No hyperbole. No exclamation. The page reads like a briefing, not a pitch.	We slide into AI-marketing voice and the CIO files us with the rest.
SOVEREIGN	Control language wins: own, operate, govern, boundary, custody, audit. Sovereignty is the worldview, not a feature.	We become 'private AI' — a category crowded with copycats, and a category the buyer does not yet trust.
SPECIFIC	We name the workload, the regulator, the failure mode. 'AESCSF', 'CPS 234', 'substation operator', 'underwriting model' — not 'regulated industries' in the abstract.	We become generic enterprise software. The CIO reads two paragraphs and does not know which problem we solve.

4. The five voice principles

Every piece of writing should be testable against these. They are ordered by how often they get violated, hardest first.

01 Lead with the risk the buyer carries, not the technology we built.

Why. The CIO's emotional driver is fear — of breach, of regulator action, of being the one who let unsanctioned AI in. The hyperscaler tagline talks about productivity. Ours has to talk about control.

DO	*AI that stays inside your trust boundary, with the audit trail your regulator expects.*
DON'T	*Powerful enterprise AI infrastructure for forward-thinking organisations.*

Test. Could you remove 'AI' from this sentence and it still sounds like risk-management language? If yes, you're close.

02 Use 'Own, Operate, Govern' as a structural device, not a slogan.

Why. The POV from the workshop is a three-verb test. When it appears in copy, it should be doing structural work — three sub-claims, three proof points, three sections — not decoration.

DO	*You own the data path. You operate the runtime. You govern who sees what, and you can prove it.*
DON'T	*Arda gives you sovereign AI you can own, operate, and govern.*

Test. If you can remove the phrase from the sentence and the sentence still works, you're using it as decoration. Cut it or rebuild around it.

03 Be specific about the world the buyer lives in.

Why. Energy networks have substations, dispatch decks, AESCSF, AEMO. Banks have CPS 234, CPS 230, claims handlers. 'Regulated industry' in the abstract sounds like every other vendor. Naming the regulator and the workflow is what makes us sound like we already work with them.

DO	*For a transmission operator running an AEMO-aligned cyber program, AI cannot mean another data path out of the boundary.*
DON'T	*Designed for organisations with strict compliance requirements.*

Test. Could the same sentence be on Microsoft, AWS, and Anthropic's websites? If yes, rewrite it for Arda.

04 Earn every claim. Footnote what you can prove. Drop what you can't.

Why. This audience reads claims as risks. '30–50% developer productivity' on the current site is unfootnoted and reads as marketing copy to a CIO who has been burnt by it before. Either back it with a customer or remove it.

DO	*In our pilot at [named operator], engineering teams shipped [X] in [Y]. (Footnote source.)*

DON'T	*Increase developer speed by 30–50%.*

Test. If the regulator or a journalist asked us to evidence this claim, do we have it? If not, it doesn't ship.

05 Engineering depth is a credential, not the headline.

Why. The strongest sales material we have (the Internal POV, First Call Deck) leads with a worldview and uses engineering as proof. The website does the opposite. The voice should follow the sales material, not the homepage.

DO	*[Hero: the worldview.] [Section 2: how it works inside the boundary.] [Section 3: the deployment options.] [Section 4: the receipts — customers, certifications, audits.]*
DON'T	*[Hero: 'The Private Enterprise AI Infrastructure Layer.']*

Test. Does the page open with what the buyer cares about, or what the engineer who built it is proud of?

5. Voice compass

Four axes. The marker on each axis is where Arda sits in writing — and where it does not. Use these when you're not sure how to phrase something.

Axis	We are / We are NOT	What it means in copy
Posture	Operator / Vendor	We talk like someone who runs the system, not someone selling it. 'You will need to think about X' beats 'we offer X'.
Register	Boardroom / Hackathon	We default to language a CIO could read in a risk committee. Engineering precision is welcome; engineering performance is not.
Posture on AI	Sceptic-builder / Believer-evangelist	We are not in the AI-is-magic camp. We are in the 'this is the only sensible way to deploy it inside a regulated boundary' camp.
Density	Compressed / Expansive	Few words, high information density. A claim, a proof, a next sentence. The current website over-bullets — short fragments without sentences. We want short sentences.
Emotion	Quiet confidence / Loud excitement	Never urgent, never breathless. We are the calm voice in a room full of AI panic. That is the product.
Authority	Earned by precision / Claimed by adjectives	'Unmatched security' is an adjective claim. 'Audit-grade logs aligned to AESCSF' is a precision claim. We do precision.

6. Lexicon — the words we own and the words we avoid

These are the working lists. Not exhaustive. They will evolve as we win deals and read the language back from real customers.

6.1 Owned vocabulary — words that earn their keep

Use these consistently. Repetition is a feature, not a bug — it's how a category gets named.

Word / phrase	Why we use it	How it shows up
Sovereign / sovereignty	Our category claim. From the workshop POV: control + operation + governance, not geography.	'Sovereign AI for regulated enterprises.' Always disambiguate first time on the page: footnote or in-line: 'sovereignty here means control, operation and governance — not location.'
Own, Operate, Govern	The three-verb POV. Buyer-grade, structural, memorable.	Three sections of the homepage. Three slides. Three lines of an email. Recurring rhythm.
Trust boundary	The CISO's word. It is the line on the architecture diagram. Using it signals we already think like them.	'Stays inside your trust boundary.' 'Custody never leaves the boundary.'
Custody (of data)	Stronger than 'access' or 'storage'. Carries the legal/regulatory connotation we want.	'You keep custody of the data path, the model weights, and the audit trail.'
Audit-grade / audit trail	Precision claim. Connects to AESCSF, CPS 234, SOC 2 — not 'logs'.	'Every prompt, every model call, every output — audit-grade.'
Regulated workload	More specific than 'enterprise data'. Names the constraint.	'For the regulated workloads that the hyperscalers can't host today.'
Inside the boundary / outside the boundary	The architectural lens that runs through the POV deck.	'Hyperscaler AI sends the prompt outside the boundary. Arda doesn't.'
Deploy / deployment	Operator language. We deploy; we don't 'roll out' or 'launch'.	'Deployable in your private cloud, your data centre, or hybrid.'
The receipts	Internal shorthand for 'evidence we can show a regulator'. Not a public phrase, but the discipline behind every claim.	Every public claim has a receipt. If it doesn't, cut it.

6.2 Banned and high-caution vocabulary

Some are banned outright. Others are conditional — usable only with proof or qualification.

BANNED — never publish	HIGH-CAUTION — only with proof
Transformative / transformational	Powerful (only with a named workload)
Cutting-edge / state-of-the-art	Productivity (only with a footnoted source)
10x / 100x / X times faster	Faster (only with the customer named)
AGI / superintelligence	Cost-effective (only with a published comparison)
Revolutionise / disrupt	Enterprise-grade (replace with the certification name where possible)
Unleash / unlock potential	Scalable (replace with the actual scale curve)
Effortless / one-click	Best-in-class (replace with the third-party benchmark)
Hyperscale (as a positive)	Innovative (replace with the specific innovation)
Ship faster, stay private (as currently used — too punchy for buyer)	Secure (replace with the specific control)
Future-proof	Trusted (replace with 'used by [named customer]')

ON 'SOVEREIGN AI'

The phrase is becoming crowded. Mistral, Oracle, CGI, Zadara and a dozen others now use it. We don't abandon it — it's our category claim — but we always carry our own definition with us. Treat 'sovereign' like a contested election: every time we use it, we restate what we mean by it. Three verbs, one boundary, audit-grade. Owned, operated, governed. The definition is the moat.

6.3 Naming and house style

- Company name is Arda. Not ARDA in body copy (use small caps in headers if the visual brand allows; otherwise 'Arda'). Use ARDA only in product/system names where it's part of the literal name (e.g. 'Arda Core', not 'ARDA Core' — keep this consistent across the site).
- Product names: Arda Core, Arda Code, Arda Index, Arda Graph. No version numbers in copy unless we are explicitly comparing versions.
- Australian English is the house standard for AU-facing content (organisation, programme, customise). US English is acceptable for US-facing content. Mixed spelling within a single page is not.
- Numerals: nine and below as words, 10+ as digits. Currency: AUD by default in AU content, USD with currency code (USD 3,000) when contextually relevant.

- Acronyms: spell out on first use, then acronym. Keep the regulator name visible — 'Australian Energy Sector Cyber Security Framework (AESCSF)' on first reference, 'AESCSF' after. The same goes for CPS 234, SOCI Act, AEMO.
- Sentence case for headlines and buttons. Title case is journalistic and outdated for B2B.
- Em dashes — used sparingly, no spaces in some house styles, but we use spaces around them — for clarity in legal and regulatory copy.
- Avoid the Oxford comma in lists of three; we are British-default in punctuation. Exception: when ambiguity demands it.

7. Reference brands

Two anchors. Both confident, both technical, both built credibility over years. Neither is a template — Arda's voice has to be Arda's. But each models a discipline we want to learn from.

7.1 Zscaler — what we borrow

What it does well. Zscaler took 'zero trust' from a security concept and made it a corporate worldview. Their copy never sells features first. It sells a way of seeing the network. Product names, customer logos and certifications come second — after the worldview is established.

What we take	Why	How
The worldview-first opening.	It works on a CIO. Lead with how the world is changing, then what they should do about it, then what we provide.	Hero is the thesis. Section 2 is the implication for the buyer. Section 3 onwards is the product.
Customer logos and certifications as proof, not decoration.	When Zscaler shows logos, they're tied to use cases. We should do the same — every logo is a story, not wallpaper.	Wait until we have first 2–3 named customers before we use a logo wall. Until then, anonymised case studies.
Compliance acronyms as currency.	FedRAMP, IL5/IL6, ISO 27001 — Zscaler treats them as proof points, not banner badges.	AESCSF, SOC 2, ISO 27001, CPS 234, SOCI alignment. Same approach: cite where it matters in the page, don't decorate the header.
The boardroom register.	Zscaler can be read aloud in a CIO/CISO meeting without sounding 'startup'.	If you can't imagine a customer reading our line aloud to their CIO, rewrite it.

What we don't take from Zscaler

- Their scale-claim posture ('largest cloud security platform') — we are early; we win on precision and trust, not size.
- Their high-volume content engine — we should publish less and prove more. The opposite of content marketing.

7.2 Mistral — what we borrow

What it does well. Mistral codified 'European sovereign AI' in a way that turned a geographic accident into a strategic identity. Their copy is sober, scientific, low on adjectives, and consistent — every page extends the same argument. They never apologise for being smaller than OpenAI; they reframe what 'better' means.

What we take	Why	How
Sovereignty as identity, not feature.	Mistral makes 'sovereign' carry the brand, not the product	Every product page opens with how that product expresses

What we take	Why	How
	page. We need the same — sovereignty is the worldview that all products inherit from.	sovereignty (Own / Operate / Govern). Not a generic product blurb.
Reframing 'better' on our own terms.	Mistral does not try to win on benchmarks alone. They win on openness, controllability, and European jurisdiction. We win on Australian sovereignty, audit-grade governance, and inside-the-boundary deployment.	Never compare ourselves to OpenAI/Anthropic on raw model quality. Compare on the dimensions where we are structurally different — control, custody, audit.
Restraint.	Mistral's homepage is short. Few claims, high density. We do the same.	Resist the urge to fill the page. If a sentence isn't earning, cut it.
Scientific posture.	Their language is precise. They say what the model does and doesn't do.	We are honest about scope. If we deploy with hyperscaler GPU as a runtime, we say so. If a workload is out of scope, we say so.

What we don't take from Mistral

- Their open-weights identity is theirs — that's their moat. Our moat is sovereign deployment + governance, not openness. We don't borrow that flag.
- Their European-policy framing — we have an Australian policy frame (SOCI, AESCSF, CPS 234) that's at least as strong. Use ours, don't echo theirs.

THE REFERENCE TRIANGLE

Zscaler gives us the boardroom register and the certification-as-proof discipline. Mistral gives us the sovereignty-as-identity stance and the restraint. Arda's voice sits between the two — closer to Zscaler in audience (the regulated CIO), closer to Mistral in posture (sovereignty as worldview, not just a feature on the page).

8. Channel guidance

The voice is one. The expression varies by channel. For each, the table below names the lead register, the typical length, the things to do more of and the things to do less.

8.1 Website and marketing copy

Channel role	What good sounds like	What bad sounds like
Homepage hero	Worldview line + qualifier. 8–14 words. Names the buyer's problem, not our product.	'The Private Enterprise AI Infrastructure Layer' (current) — speaks to engineers, not the CIO.
Sub-hero	One sentence on what we do, framed in the buyer's language: control, custody, audit, regulator alignment.	Bullet list of features.
Solution / product page	Open with the workload, not the product. 'For [workload] in [sector]' — then how Arda works inside it.	Generic product description that could be on any AI vendor's site.
Industry pages (Energy, FSI, etc.)	Name the regulator, the workflow, the failure mode. Show we already think in their world.	'For organisations in regulated industries.'
Footer / trust band	Certifications + named customer logos (when available). No badge collage without context.	Decorative badge wall.
Long-form / POV pieces	Strategic essays that extend the worldview — like the Internal POV deck made public. Argue, don't pitch.	Listicles. SEO bait. Trend posts.

8.2 Sales materials and decks

Asset	What good sounds like	What bad sounds like
First-call deck	Worldview > the gap > how Arda fills it > what it means for them > next step. Six to eight slides max.	Twenty slides of feature lists.
Proposal	Customer's problem in their words on page 1, our framing on page 2, scope on page 3, commercials at the back.	Boilerplate template that could be sent to anyone.
Architecture / security pages	Diagram first, words second. Plain English on the diagram caption. Engineering depth in the	Walls of text without a diagram. Or: a diagram with no caption.

Asset	What good sounds like	What bad sounds like
	appendix.	
Email follow-up	One reference back to the conversation, one new piece of evidence, one specific next step. Three short paragraphs.	'Just circling back' / 'Thought I'd reach out'.

8.3 Founder and exec voice — LinkedIn, PR, talks

Erik and Paul (and increasingly Alex, Stefan, Costin) are public-facing. Same DNA, slightly warmer register — first person allowed, lived examples encouraged, but the discipline holds: sober, sovereign, specific.

Format	What good sounds like	Cadence guideline
LinkedIn post (Erik / Paul)	A specific thing they saw this week + what it tells them about regulated AI adoption. 80–180 words. No threadwormy hooks.	1–2 a week per founder, max.
LinkedIn comment	A real opinion, expressed politely, on someone else's post. Not 'great post!' Not 'fully agree, our customers see this too.'	Daily-ish if it's a real reaction, never if it's not.
Op-ed / long form	A worldview piece tied to a current event (regulator announcement, breach, AI policy moment). 800–1,200 words. Argumentative.	Quarterly, in a publication the buyer reads — AFR, InnovationAus, banking/energy trade press.
Conference talk / podcast	A clear thesis, three substantiations, a sharp closing line. No 'and we're hiring' at the end.	Pick the venues where the buyer already is. Don't optimise for AI conferences.

FOUNDER VOICE — THE DISCIPLINE

Erik and Paul are the brand. Anything they post is read by the next CIO we want to win. The cost of being too quiet is low. The cost of being too loud, off-thesis, or unfootnoted is high. Default to less.

8.4 Customer and partner comms

Touchpoint	What good sounds like	Trap to avoid
Pilot kick-off email	Short. Names the success criteria for the pilot. Sets a calendar cadence. Plain English.	Marketing language in a pilot email. They already bought; stop selling.
Pilot status update	What we did, what worked, what didn't, what's next. Honest about the 'didn't'.	Selective reporting. CIOs see through it and lose trust faster than for any other failure.

Touchpoint	What good sounds like	Trap to avoid
Issue / incident comms	Acknowledge, scope, root cause, mitigation, prevention. In that order. Plain English, no hedging.	Legal-influenced passive voice ('an issue was identified').
Partner co-marketing (Oracle, others)	Make the partnership about the customer's gain, not the logo trade.	Joint logos with no joint substance.
Customer reference / case study	Named customer (with permission), named workload, named outcome with numbers we can defend.	Anonymised quotes that sound made up.

9. Sample copy

Working examples. Treat them as starting points for the next round of writing — not finished copy. Each one shows the principles in motion.

9.1 Homepage hero — three options to test

OPTION A — BUYER'S RISK LENS (RECOMMENDED)

AI you can defend to your regulator.

Arda is the sovereign AI platform regulated enterprises own, operate, and govern — inside their own trust boundary, with the audit trail their CISO and their regulator expect.

OPTION B — CATEGORY CLAIM

The AI you own. Operate. Govern.

For Australia's most regulated organisations — energy networks, banks, insurers — Arda runs the AI workloads that cannot leave your boundary.

OPTION C — OPERATOR'S VOICE

Sovereign AI for the workloads that can't leave the boundary.

Chat, code and search inside your environment. Your data stays inside. Your model decisions stay yours. Your audit trail is ready when the regulator asks.

9.2 Elevator — one sentence

'**Arda is the sovereign AI platform regulated enterprises own, operate and govern — so they can adopt AI without changing their security posture.**'

Decision pending: drop 'security leader' from the working sentence in the synthesis (Simon's note). Above is the revision — buyer-led, three-verb structure, ends on the buyer's promise.

9.3 LinkedIn post — Erik (sample)

SAMPLE — SHORT, SPECIFIC, ONE SUBSTANTIATION

A CIO at one of Australia's largest distributors told me last week:

'I can't be the executive who let unsanctioned AI into our control room.'

That sentence is the regulated AI market in 12 words.

It's not anti-AI. It's anti-leakage. The same CIO has 200 engineers asking for Copilot, and a board asking for 'an AI plan'. The answer can't be 'no'. It also can't be 'yes, send it all to a US frontier lab'.

Sovereign AI — the boring kind, that runs inside the boundary, with audit trails — is what they're actually buying.

Why it works: opens with a specific buyer; quotes verbatim; reframes the choice; closes on our category in three short clauses. ~110 words.

9.4 Cold outbound email — first touch

SUBJECT: AESCSF-ALIGNED AI FOR [THEIR ORGANISATION]

[Name],

I work with three Australian energy operators on the same problem: the board wants an AI plan, the CISO won't sign off on hyperscaler tooling for anything that touches operational data, and the team is running out of patience.

Arda runs the AI workloads — chat, coding, search over internal data — inside your own trust boundary. No prompts leave. No training on your data. Audit-grade logs aligned to AESCSF and SOCI.

Worth 25 minutes? I can show you the architecture and how a transmission peer of yours is using it.

Paul Coghlan — Arda Technologies

Why it works: names the regulator and sector; reflects three pains the CIO actually has; specific architectural claim; specific ask; signals peer adoption. Adapt the pain framing for FSI.

9.5 Pilot follow-up — week 2 update

SUBJECT: [CUSTOMER] X ARDA — WEEK 2 UPDATE

Three lines on where we are.

Working: chat is live in your environment for 28 users. Median response 1.4s. No prompts logged outside the boundary — happy to walk your CISO through the audit log.

Not yet working: Confluence connector is two days behind plan — we hit an auth quirk in your SSO config. Workaround in place; permanent fix Friday.

Next: Stefan and your data team are scoping the second use case (claims handler assist) for week 4.

Same time next week?

Why it works: working / not working / next — the discipline. Names the failure before they have to ask. Three short blocks; CIO can read in 30 seconds and forward to their team.

10. Boilerplates

Fixed building blocks. Use verbatim — copy/paste — until we've earned the right to revise them with customer evidence.

10.1 The 25-word boilerplate

USE IN: FOOTER, PARTNER DIRECTORIES, PRESS BOILERPLATE, AWARDS SUBMISSIONS

Arda is the sovereign AI platform that regulated enterprises own, operate and govern — so they can adopt AI without changing their security posture.

10.2 The 50-word boilerplate

USE IN: ABOUT PAGE, PITCH DECK CLOSING SLIDE, BOARD PAPER BACKGROUND

Arda is the sovereign AI platform that regulated enterprises own, operate and govern.

Built for Australia's most regulated organisations — energy networks, banks, insurers — Arda runs AI workloads inside the customer's trust boundary, with the audit trail their CISO and their regulator expect. Chat, coding, and enterprise search — without sending a prompt or a line of code outside.

10.3 The proof line

A single repeatable line that bundles every claim into one sentence — to use as a closing line, a tagline alternative, or a tweet.

USE IN: CLOSING LINES, SOCIAL BIOS, SIGNATURE BLOCKS, TAGLINE TESTS

Inside the boundary. Audit-grade. Sovereign by design.

11. Before and after — current website

Selected lines from ardacomputing.com (current live site) rewritten against the principles. These are illustrative of direction, not the final web copy. The full rewrite belongs in the website rebuild brief — see Section 14.

Hero

BEFORE — current site	AFTER — on-voice
The Private Enterprise AI Infrastructure Layer *Arda is secure AI Infrastructure.* *Keep data private, enhance productivity, reduce cost and increase visibility — all with enterprise level security.*	AI you can defend to your regulator. Arda is the sovereign AI platform regulated enterprises own, operate, and govern — inside their own trust boundary, with the audit trail their CISO and regulator expect.

Why: Lead changes from engineering pride ('infrastructure layer') to the buyer's risk frame. Three-verb POV becomes structural. Drops productivity from the hero — it's a side-effect, not the promise.

Solutions header

BEFORE — current site	AFTER — on-voice
Our Solutions *Deployable in SaaS, air-gapped, or classified environments.*	Three workloads, all inside your boundary Chat, code, and enterprise search — deployable in your private cloud, your data centre, or hybrid.

Why: 'Solutions' is generic. Naming the three workloads upfront does the work. Deployment language stays operator-grade.

Cost section

BEFORE — current site	AFTER — on-voice
Cost effective *Your custom infrastructure, deployed and managed by us.* *Throttle or expand service at any time, set operating parameters to deploy AI efficiently.*	Costs that fall as usage grows. Open-weight models on dedicated infrastructure. Costs are flat once deployed — usage compounds inside the same envelope, instead of multiplying with every API call.

Why: 'Cost effective' is an adjective claim — replace with the actual cost curve. Borrow the strongest line from the First Call Deck (3-year cost comparison) and put it forward.

Productivity claim

BEFORE — current site	AFTER — on-voice
Increase developer speed by 30–50%.	Cut the time engineers spend hunting through your codebase — without sending a line of code outside your boundary.

Why: 30–50% is unfootnoted. Replace with the qualitative claim until we have a named pilot to cite, then re-introduce the number with a footnote.

Industries

BEFORE — current site	AFTER — on-voice
Regulated Utilities *Healthcare* *Banking* *Telecommunications*	Energy networks (T&D) Banks, insurers and super funds (FSI) Critical infrastructure operators under SOCI

Why: Match the segment commitment from Phase 1 — energy primary, FSI secondary, broader CI as adjacency. Drops healthcare and telco from the front page until they are real.

Closing CTA

BEFORE — current site	AFTER — on-voice
Ship faster. Stay private.	Adopt AI without changing your security posture.

Why: 'Ship faster, stay private' is too punchy for the buyer (it sounds like a developer-tools tagline). The replacement is plainer and lands directly on the CIO's brief.

12. Quick reference card

One page. Print it. Pin it next to whoever is writing.

THE 60-SECOND CHECK — BEFORE YOU PUBLISH

1. Could a CIO read this aloud in their risk committee without being embarrassed?
2. Have you led with risk, control or governance — not productivity or cost?
3. Is at least one sentence specific to the buyer's world (regulator, workload, sector)?
4. Does every claim have a receipt — a customer, a benchmark, a certification, a published source?
5. Is there a banned word in the draft? (transformative, cutting-edge, 10x, future-proof, unleash, revolutionise)
6. Is the page shorter than the previous version? If not, are the extra words earning their place?
7. Could this sentence sit on AWS, Microsoft and Anthropic's site without changes? If yes — rewrite for Arda.

Owned words — keep using

Sovereign / sovereignty • Own, Operate, Govern • **Trust boundary** • Custody (of data) • **Audit-grade / audit trail** • Regulated workload • **Inside the boundary** • Deploy / deployment • **Receipts (internal discipline)**

Banned words — never publish

transformative • cutting-edge • 10x / 100x • AGI • revolutionise • disrupt • effortless • one-click • future-proof • unleash • unlock potential

13. Open questions and what to validate

This v1 makes commitments where the strategy is clear and flags where it isn't. Three things to validate before this document goes to v2.

Open question	What we've assumed in v1	How to resolve
Deployment posture (on-prem vs hyperscaler-only) — unresolved per Phase 1 synthesis.	Voice talks about deployment as flexible: 'private cloud, on-prem, or hybrid'. Avoids picking a side.	Once Erik / Paul / Alex land the call, update Sections 6, 8 and 9 to reflect the chosen language. If hyperscaler-only wins, drop 'on-prem' and 'data centre' references and lead with 'private cloud / VPC inside your hyperscaler account'.
First named customer — when can we cite one?	Sample copy uses placeholders ('a transmission peer'). Productivity claims are softened until we can cite a real one.	First two pilots concluded → revisit Section 11 (Before/After) and Section 9 (Sample copy) with named customer + named outcomes.
Pricing model and language — Phase 2.	v1 avoids price language entirely. Boilerplate uses no numbers.	When pricing model is set, add a Section 6.3 entry on how we talk about price (per workload? per environment? per organisation?). Banned/cautioned word list to be revised in light of it.
Australian English vs international.	House style is Australian English. International audiences served by separate localisation.	If the AEMO/energy beachhead extends to APAC adjacencies (NZ, Singapore), confirm whether AU spelling stays or shifts to international English.
Founder posture for public speaking.	Section 8.3 sketches the discipline; not yet pressure-tested with Erik / Paul.	30-minute working session with Erik and Paul — react to the sample LinkedIn post (§9.3) and adjust the cadence guidance.

14. Hand-off into the next stage

Per the GTM plan (Phase 1 synthesis, Section 3 of the handoff doc), the next operative use of this document is the website rebuild — Weeks 5–8 of the 90-day plan, ICP messaging-and-positioning swimlane.

14.1 What this document hands over

- Voice DNA, principles and lexicon — the rules anyone writing on Arda's behalf is now expected to follow.
- Channel guidance for web, sales, exec voice and customer comms — so the same voice runs across surfaces.
- Sample hero and sub-hero options for the homepage — three to test in the rebuild.
- Before/after rewrites of selected current site copy — directional, not final.
- Boilerplates (25-word, 50-word, proof line) — usable verbatim from today.

14.2 What still needs doing

- Full website rebuild brief — content architecture, page-by-page copy, IA decisions. This document feeds it; it does not replace it.
- Resolve the on-prem vs hyperscaler-only question. The voice can ship without it; the website rebuild cannot.
- Visual brand pass — colour, typography, photography. Voice is one half of the brand; visual is the other. They should be developed against each other, not in isolation.
- Stakeholder review — Erik, Paul, Alex on the principles; Stefan on the engineering language (so we're not undercooking what's technically true).

14.3 Stage move — when this work shifts platform

This document is the strategy and rules. The next stage — building the new ardacomputing.com — moves from Cowork (strategy / docs) to Claude Code in VS Code (build). The hand-off prompt to use when ready:

COPY / PASTE — CLAUDE CODE PROMPT FOR THE WEBSITE BUILD

```
I'm rebuilding ardacomputing.com. The brand strategy and tone of voice
are locked — see ARDA_Tone_of_Voice_v1_2026-05-01.docx and
ARDA_Phase1_Workshop_Synthesis_2026-04-27.docx.

Read both documents end-to-end before writing anything. Then propose:

1) site information architecture (homepage + product pages + industry
pages + about + resources),

2) recommended hero (pick one of the three options in §9.1, justify your
choice),

3) page-by-page copy outline applying the tone of voice (use the
Before/After in §11 as the directional reference),
```

```
4) tech stack recommendation (assume we want the site fast, statically
generated, and easy for a non-engineer to update copy on).

Constraints: Australian English. No banned words from §6.2. Every claim
must have a receipt or be cut. Lead buyer = CIO. Primary segment = AU
energy networks, secondary = FSI.

Push back on me on anything in the ToV that does not map cleanly to a
buildable site.
```

Paste this into Claude Code in VS Code with both files attached when you're ready to move into the build stage. The prompt asks for an IA + copy plan first, not code — you stay in driver's seat for the structural decisions.

Appendix — source material

- Phase 1 Workshop Synthesis (27 April 2026) — *ARDA_Phase1_Workshop_Synthesis_2026-04-27.docx*
- Phase 1 Workshop Handoff (April 2026) — *ARDA_Phase1_Workshop_Handoff.md*
- Internal POV — *arda_pov_regulated_ai_adoption.pdf*
- First Call Deck — *ARDA_First_Call_Deck.pdf*
- Executive Slides — *Arda Executive Slides.pdf*
- Live website (May 2026 capture) — *ardacomputing.com*
- Reference brand — *Zscaler.com, Mistral.ai (May 2026)*

Arda Technologies — Tone of Voice v1 — 1 May 2026 — Prepared by Simon Bowker, Hepburn Advisory.